

Week 7: Nonlinear models

Slava Mikhaylov

PUBLG088 Advanced Quantitative Methods

Week 7 Outline

Moving Beyond Linearity

Polynomial Regression

Step Functions

Splines

Local Regression

Generalized Additive Models

Moving Beyond Linearity



Linearity and social reality

- ▶ Social world is almost never linear.
- ▶ Often the linearity assumption is good enough.
- ▶ When linearity doesn't hold we can use
 - ▶ polynomials,
 - ▶ step functions,
 - ▶ splines,
 - ▶ local regression, and
 - ▶ generalized additive models
- ▶ These models offer a lot of flexibility, without losing the ease and interpretability of linear models.

Polynomial regression

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \cdots + \beta_d x_i^d + \epsilon_i$$

Degree-4 Polynomial



Details

- ▶ Create new variables $X_1 = X$, $X_2 = X^2$, etc and then treat as multiple linear regression.
- ▶ Not really interested in the coefficients; more interested in the fitted function values at any value x_0 :

$$\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0 + \hat{\beta}_2 x_0^2 + \hat{\beta}_3 x_0^3 + \hat{\beta}_4 x_0^4.$$

- ▶ Since $\hat{f}(x_0)$ is a linear function of the $\hat{\beta}_\ell$, can get a simple expression for **pointwise-variances** $\text{Var}[\hat{f}(x_0)]$ at any value x_0 .
- ▶ In the figure we have computed the fit and pointwise standard errors on a grid of values for x_0 . We show $\hat{f}(x_0) \pm 2 \cdot \text{se}[\hat{f}(x_0)]$.
- ▶ We either fix the degree d at some reasonably low value, else use cross-validation to choose d .

Details continued

- ▶ Logistic regression follows naturally. For example, in figure we model

$$Pr(y_i > 250|x_i) = \frac{\exp(\beta_0 + \beta_1 x_i + \beta_2 x_i^2 \cdots + \beta_d x_i^d)}{1 + \exp(\beta_0 + \beta_1 x_i + \beta_2 x_i^2 \cdots + \beta_d x_i^d)}.$$


- ▶ To get confidence intervals, compute upper and lower bounds on **the logit scale**, and then invert to get on probability scale.
- ▶ Can do separately on several variables just stack the variables into one matrix, and separate out the pieces afterwards (see GAMs later).
- ▶ Caveat: polynomials have notorious tail behavior – very bad for extrapolation.
- ▶ Can fit using $y \sim \text{poly}(x, \text{degree} = 3)$ in formula.

Step Functions

Another way of creating transformations of a variable – cut the variable into distinct regions.

$$C_1(X) = I(X < 35), C_2(X) = I(35 \leq X < 65), \dots, C_3(X) = I(X \geq 65)$$

Piecewise Constant



Step functions continued

- ▶ Easy to work with. Creates a series of dummy variables representing each group.
- ▶ Useful way of creating interactions that are easy to interpret. For example, interaction effect of *Year* and *Age*:

$$I(\text{Year} < 2005) \cdot \text{Age}, I(\text{Year} \geq 2005) \cdot \text{Age}$$

would allow for different linear functions in each age category.

- ▶ In R: $I(\text{year} < 2005)$ or $\text{cut}(\text{age}, c(18, 25, 40, 65, 90))$.
- ▶ Choice of cutpoints or **knots** can be problematic. For creating nonlinearities, smoother alternatives such as **splines** are available.


Piecewise polynomials

- ▶ Instead of a single polynomial in X over its whole domain, we can rather use different polynomials in regions defined by knots. E.g. (see figure)


$$y_i = \begin{cases} \beta_{01} + \beta_{11}x_i + \beta_{21}x_i^2 + \beta_{31}x_i^3 + \epsilon_i & \text{if } x_i < c; \\ \beta_{02} + \beta_{12}x_i + \beta_{22}x_i^2 + \beta_{32}x_i^3 + \epsilon_i & \text{if } x_i \geq c. \end{cases}$$

- ▶ Better to add constraints to the polynomials, e.g. continuity.
- ▶ **Splines** have the “maximum” amount of continuity.


Piecewise Cubic




Continuous Piecewise Cubic



Cubic Spline



Linear Spline



Linear Splines

- ▶ A linear spline with knots at ξ_k , $k = 1, \dots, K$ is a piecewise linear polynomial continuous at each knot.
- ▶ We can represent this model as

$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+3} b_{K+3}(x_i) + \epsilon_i,$$

where the b_k are **basis functions**.

$$b_1(x_i) = x_i$$

$$b_{k+1}(x_i) = (x_i - \xi_k)_+, \quad k = 1, \dots, K$$

Here the $(\)_+$ means **positive part**; i.e.

$$(x_i - \xi_k)_+ = \begin{cases} x_i - \xi_k & \text{if } x_i > \xi_k; \\ 0 & \text{otherwise.} \end{cases}$$

Cubic splines

- ▶ A cubic spline with knots at ξ_k , $k = 1, \dots, K$ is a piecewise cubic polynomial with continuous derivatives up to order 2 at each knot.
- ▶ We can represent this model with truncated power basis functions

$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+3} b_{K+3}(x_i) + \epsilon_i,$$

where the b_k are **basis functions**.

$$b_1(x_i) = x_i; b_2(x_i) = x_i^2; b_3(x_i) = x_i^3;$$


$$b_{k+3}(x_i) = (x_i - \xi_k)_+^3, \quad k = 1, \dots, K$$

where

$$(x_i - \xi_k)_+^3 = \begin{cases} (x_i - \xi_k)^3 & \text{if } x_i > \xi_k; \\ 0 & \text{otherwise.} \end{cases}$$



Natural cubic splines

- ▶ A natural cubic spline extrapolates linearly beyond the boundary knots.
- ▶ This adds $4 = 2 \times 2$ extra constraints, and allows us to put more internal knots for the same degrees of freedom as a regular cubic spline.




- Fitting splines in R is easy: $bs(x, \dots)$ for any degree splines, and $ns(x, \dots)$ for natural cubic splines, in package *splines*.

Natural Cubic Spline



Knot placement

- ▶ One strategy is to decide K , the number of knots, and then place them at appropriate quantiles of the observed X .
- ▶ A cubic spline with K knots has $K + 4$ parameters or degrees of freedom.
- ▶ A natural spline with K knots has K degrees of freedom.
- ▶ On the figure, comparison of a degree-14 polynomial and a natural cubic spline, each with 15df: $ns(\text{age}, df = 14)$ and $poly(\text{age}, deg = 14)$.



Smoothing splines

- ▶ Consider this criterion for fitting a smooth function $g(x)$ to some data:

$$\underbrace{\text{minimize}}_{g \in S} \sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt$$

- ▶ The first term is RSS, and tries to make $g(x)$ match the data at each x_i .
- ▶ The second term is a **roughness penalty** and controls how wiggly $g(x)$ is. It is modulated by the **tuning parameter** $\lambda \geq 0$.
- ▶ The smaller λ , the more wiggly the function, eventually interpolating y_i when $\lambda = 0$.
- ▶ As $\lambda \rightarrow \infty$, the function $g(x)$ becomes linear.

Smoothing splines continued

- ▶ The solution is a natural cubic spline, with a knot at every unique value of x_i .
- ▶ The roughness penalty still controls the roughness via λ .

Smoothing splines details

- ▶ Smoothing splines avoid the knot-selection issue, leaving a single λ to be chosen.
- ▶ The algorithmic details are too complex to describe here. In R, the function `smooth.spline()` will fit a smoothing spline.
- ▶ The vector of n fitted values can be written as $\hat{\mathbf{g}}_\lambda = \mathbf{S}_\lambda \mathbf{y}$, where \mathbf{S}_λ is a $n \times n$ matrix (determined by the x_i and λ).
- ▶ The **effective degrees of freedom** are given by

$$df_\lambda = \sum_{i=1}^n \{\mathbf{S}_\lambda\}_{ii}.$$


Smoothing splines – choosing λ

- ▶ We can specify df rather than λ . In R:
smooth.spline(age, wage, $df = 10$).
- ▶ The LOOCV error is given by

$$\text{RSS}_{cv}(\lambda) = \sum_{i=1}^n (y_i - \hat{g}_{\lambda}^{-i}(x_i))^2 = \sum_{i=1}^n \left[\frac{y_i - \hat{g}_{\lambda}(x_i)}{1 - \{\mathbf{S}_{\lambda}\}_{ii}} \right]^2.$$


In R: *smooth.spline*(age, wage)

Smoothing Spline



Local regression

Local Regression




- ▶ With a sliding weight function, we fit separate linear fits over the range of X by weighted least squares.
- ▶ See text for more details, and `loess()` function in R.

Generalized Additive Models

Allows for flexible nonlinearities in several variables, but retains the additive structure of linear models.

$$y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}) + \epsilon_i.$$




GAM details

- ▶ Can fit a GAM simply using, e.g. natural splines:
 $lm(wage \sim ns(year, df = 5) + ns(age, df = 5) + education)$
- ▶ Coefficients not that interesting; fitted functions are. The previous plot was produced using *plot.gam*.
- ▶ Can mix terms – some linear, some nonlinear – and use *anova()* to compare models.
- ▶ Can use smoothing splines or local regression as well:
 $gam(wage \sim s(year, df = 5) + lo(age, span = .5) + education)$
- ▶ GAMs are additive, although low-order interactions can be included in a natural way using, e.g. bivariate smoothers or interactions of the form $ns(age, df = 5) : ns(year, df = 5)$.

GAMs for classification

$$\log \left(\frac{p(X)}{1 - p(X)} \right) = \beta_0 + f_1(X_1) + f_2(X_2) + \cdots + f_p(X_p).$$



$\text{gam}(I(\text{wage} > 250) \sim \text{year} + s(\text{age}, df = 5) + \text{education}, \text{family} = \text{binomial})$